

LET’S TACKLE SYCOPHANCY FROM THE SOURCE & WHERE IT MATTERS MOST

Nowe Moore
March - April 2026

Abstract: As models become more capable, sycophancy becomes a larger safety concern through creating distrustworthy systems, propagating misinformation, and inhibiting prosocial behaviour amongst human users. Although it could contribute to an understanding of where sycophancy comes from as well as what areas of society are most exposed to its effects, scientists have yet not yet modelled how sycophancy patterns vary across context domains of human-AI interactions. The present study tasks 4 SOTA models with 640 tasks across 32 context domains represented in open-source post-training datasets to observe how sycophancy levels vary between responses to objectively disputeable and objectively falsifiable tasks, affirmatively- and negatively-phrased queries, and first- and third-person queries *across context domains*.

(Preliminary) Key Findings: (1) This study provides empirical evidence for a previously suspected objectivity effect: sycophancy is significantly higher for debateable than for objectively falsifiable queries regardless of topic. (2) Context domains stereotypically viewed as more debateable (and therefore expected to demonstrate higher sycophancy levels) actually showed below-average sycophancy levels when the objectivity of a statement was controlled for. (3) There are significant differences between not only general sycophancy levels but also relative sycophancy levels between evaluated models across topics. (4) Previously described third-person effect, where models show higher levels of sycophancy in response to first-person queries than to third-person queries, and the polarity effect, where models show higher sycophancy on affirmatively-phrased queries, did not reproduce consistently across conditions.

INTRODUCTION. Sycophancy, the continued tendency of models to tailor their responses to the user’s perceived beliefs and preferences, sometimes even at the cost of accurate, truthful information, is a considerable obstacle to what many of us imagine under robust, reliable, authentic systems of the future. While this at best obsequious and possibly even disingenuous flattery underlying models’ reactions to the surrounding world is often viewed as a capability defect, it has become increasingly more evident that sycophancy is, first and foremost, a major safety issue in need of context-sensitive control. In that spirit, sycophancy research has come a long way, though mostly treating context domain as a proxy for other variables. This study asks,

for the first time, what context domain itself contributes to sycophantic behaviour and how it could help us understand where sycophancy comes from and how to target mitigations toward areas with the highest exposure to its effects.

LIT REVIEW. Sycophancy & Safety. A sycophantic model is a model agnostic to the truth, unfaithful to its own principles, and thus incapable of performing tasks with integrity. This framing makes it obvious why sycophancy is an obstacle to capability, but it is also worth noting that sycophancy is a non-negligible safety concern.

For one, sycophancy is not invisible to human users. If and when they notice increased levels of

unwarranted flattery in the midst of their interaction with a model, they become immediately more hesitant to trust its responses[4][5]. As deployment scales, human societies will find themselves not with the helpful assistants they were promised, but with untrustworthy imposters already embedded in daily life. This can lead to a chronic, society-wide sense of doubt, eroded confidence in information, and a society that has learned to expect flattery over truth. Whether a society conditioned to distrust can sustainably survive and flourish is left as an exercise to the reader.

The outstanding reasons for why sycophancy poses a risk to AI safety are less hypothetical: it could propagate misinformation amongst unsuspecting users[4][6][8] and steer user behaviour in ways damaging to social relations[5]. The harms of misinformation are intuitive: from erosion of informed decision-making to more large-scale manipulation of public opinion. At the behavioural level, the damage is subtler but no less powerful: for example, users interacting with sycophantic models were more likely to confirm their own rightness in an inter-personal conflict and ended up less likely to entertain engaging in repair actions[5].

Perhaps the most significant safety concern regarding sycophancy is that we do not exactly know how to control it: part because we have yet to determine what properties of training data cause it and part because we know very little about when to expect deviating sycophancy in deployment. To make progress on both fronts (i.e. towards understanding the origin of sycophancy and correctly predicting it in deployment), having a better idea of what triggers sycophancy[7] would be a productive starting point. Triggers, once identified, would narrow the training data search and help surface likely causes, as well as provide empirical grounding for targeted interventions. This study attempts to make a small step in that direction.

Origins of Sycophancy. Previous research demonstrated that RLHF is likely the core am-

plifier of sycophantic behaviour[11][14][15]. At first, by observing that humans systemically select more sycophantic responses when presented with preferential data, at times at the expense of factual integrity[15]. Some two years later, researchers were also able to find a mechanistic explanation for such behaviour, by showing a link between a learnt reward and bias in RLHF data[14].

The problem with RLHF is that it is probably only powerful enough to amplify—but not introduce—sycophancy in SOTA models. The cause of sycophancy is likely rooted more deeply in the training data. This is evidenced by the fact that scaling alone is enough to amplify sycophancy (incl. on tasks where the model otherwise knows the correct answer)[18]. In other words, whatever feature causes sycophancy, more data encodes more of it. The interesting question to ask is whether we can pinpoint what types of data propagate this data more than other. Though public access to insights about the contents of training data for SOTA models is limited, trigger patterns can serve as informative proxies, and this study builds off of that.

Triggers of Sycophancy. Thanks to efforts from many different directions, we now have quite a wealth of examples of features that systemically trigger sycophancy in model responses. One of the most graspable findings was, for example, the increase of sycophancy in response to first-person—but not third-person—queries[7][17]. This shows that, for whatever reason, models weigh in the mental model of the user more compared to everyone else’s. These findings, alongside others known to the author, inform the experimental design of this study.

Perhaps more importantly for the purposes of this study, there are several more abstract features predictive of sycophantic responses. For example, an increasing context length of interaction seems to trigger increasingly more sycophantic responses from models, and that on both user and synthetic, non-user contexts[9]. The

mechanism behind how real-world interaction contexts enable this phenomenon remains unknown.

Another relevant observation is the effect of how objective versus subjective a query is. While researchers have observed that the objectivity of a statement triggers more sycophancy, researchers conducted in this area tend to use context domain as a proxy to objectivity (e.g. maths queries = objective tasks, social queries = subjective tasks), without measuring the impact of the objectivity effect in isolation[13]. That said, other research found non-negligible sycophancy levels in ‘objective’ context domains as well[12], and understanding how much query objectivity (regardless of context domain) matters may help reconcile these conflicting observations. This study offers partial resolution to that.

Sycophancy as a Function of Topic. In summary, studying sycophancy in the context of domain context has multiple benefits. First, testing the objectivity effect across domains offers a unified basis for the fragmented claims of prior work. Second, domain-level patterns in sycophancy could act as informative proxies for the kinds of training data that introduce it in the first place. Finally, some research additionally argues that sycophancy may be a feature as well as a bug and context-sensitive mitigations may be more helpful than blanket ones[2].

RESEARCH QUESTIONS. With that in mind, this study asks the following questions: (RQ1) Does the objectivity effect trigger for sycophancy hold across topics? (RQ2) Do models show consistently show above- or below-average sycophancy on some topics over others? (RQ3) Does the representation of topic in RLHF data represent the sycophancy levels for that topic?

METHODOLOGY. Topics. It all begins with the operationalisation of context domains. For this study, we sampled 1,000 (out of 17,000) preference samples from the **TULU3** wildchat datasets and prompted GPT-4o[1] to return a word or a short phrase to describe the context domain of the response. The model was explicitly instructed not to generate tags that describe a particular skill (e.g. writing or programming) but rather a more overarching topic (e.g. politics, health, art etc.). The 1,000 yielded 232 unique tags, which were subsequently manually grouped into 32 unique tags. The final tags, alongside the original tags grouped under each, are all listed in **Appendix A**.

These 32 tags were subsequently used to tag another 5,000 entries to better model the distribution of the selected tags across the whole dataset. 32 out of the 5,000 tags did not fit any of the pre-defined categories and were removed from the analysis. The final list of topic distribution in the dataset is shown in **Figure 1**.




Figure 1: Count of topics in the dataset of 5,000 entries

Tasks. The next step consisted of defining appropriate tasks to evaluate sycophancy on both objective (experimental) and subjective (baseline) queries. To do so, we composed two task datasets, ‘objective tasks’ and ‘subjective tasks’, consisting of 320 statements each (10 for every topic). The ‘objective tasks’ dataset was composed of statements about the world designed by the following requirements: (i) debatable from either perspective (i.e. agree and disagree), (ii) minimising explicit polarisation, (iii) focusing criticism on ideas, not individuals or groups, (iv) avoiding direct comparisons where possible, (v) not factually grounded (i.e. must be an opinion, not a fact).

Each statement was subsequently embedded in 2 different polarity contexts using the affirmative-polarity prefix ‘I think that’ and negative-polarity prefix ‘I don’t think that’. Additionally, we also test equivalent statements with third-person prefixes ‘They think’ and ‘They don’t think’, to in order to reproduce extend on the findings of previous sycophancy research[7][17].

The same process was followed for the ‘subjective tasks’ dataset, except with the statement requirements adjusted as follows: (i) non-debateable, verifiably incorrect (either by common sense or expert opinion), (ii) focusing on concepts and mechanisms, (iii) not opinions or

value judgements plausible-sounding enough that the error may not be immediately obvious, (iv) not grounded in definitions alone when possible (i.e. the falsehood should stem from a factual claim, not a semantic one).

Evaluation Process. An evaluation pipeline was built using UK AISI Inspect library. A model judge was asked to evaluate every response to each of the 320 topics based on a previously established rubric[5]. This rubric was selected for its direct relevance to the safety concerns motivating this study: it was developed in the context of the most relevant effects of sycophancy, including increased dependency and inhibited pro-social behaviour, making it well-suited to the safety framing of this study. The rubric went as follows: `Given the statement ‘statement’, assess to what extent the response is affirming of the statement. Grade the response on a scale 0 to 3,`

`where: 0 = non-affirming, 1 = neutral, 2 = implicit affirming, 3 = explicit affirming.`

Judge Validation. 64 responses (1 from each topic on each task) were sampled across the collected data and presented to human evaluators trained on to evaluate model responses according to the same rubric: by thoroughly examining the model response to the task and grading the response on a scale 0-3 based on the above guidelines. Cohen’s Kappa was used to measure agreement between model and human evaluators, confirming that the model judge was reliable enough to draw statistically meaningful conclusions from.

Models. The present study evaluates 4 SOTA models: Claude Sonnet 4.6, Grok 4 Fast, GPT-5.4, and Gemini 3 Flash Preview. All models responded to tasks in 3 epochs to mitigate the effect of randomness.

Author’s Note: While the above section describes the full methodology in detail, this version only processes partial results: 2 out of 4 models have so far been evaluated on 12 out of 32 topics (representing top, middle, and bottom end of the distribution in Figure 1). Model performance on the remaining 20 tasks alongside judge validation data is due to be collected in the upcoming weeks, following the receipt of a recently approved Small Rapid Grant. Conclusions are deferred to the final version of this paper, where the complete dataset will allow for more robust interpretation.

RESULTS & HOW TO READ THEM. Due to the goals of this study and a relatively small dataset (only 10 tasks per topic), the main value of the collected data is in the comparisons: subjective versus objective tasks, affirmative versus negative queries, first versus third person queries, and between individual models and individual topics.

Subjective versus Objective Queries. Figures 2, 3, 4, 5 show how the evaluated models tested on debateable, subjective tasks. Recall that these tasks involved clearly stating the user’s position on a statement with no objectively correct answer (e.g. ‘I think X’ or ‘I don’t think X’) and asking the model for an opinion. Figures 6, 7, 8, 9 show how models performed on easily disputeable tasks. Objective queries serve

as a baseline for the level of sycophancy that models show in response to queries where there is a clear, objectively correct answer. Subjective queries, on the other hand, are expected to trigger some level of sycophancy, as there is no objectively correct answer and the model may be more likely to cater to the user’s perceived beliefs and preferences. The difference in model levels of sycophancy between the two query types represents empirical evidence for the objectivity effect. The interesting question is whether this effect holds across topics, and if so, how much of it can be explained by the representation of the topic in RLHF data.

Affirmative versus Negative Queries. In all figures, affirmative queries (i.e. using prefix

‘I think’ or ‘They think’) are colour coded with salmon while negative queries (i.e. using prefix ‘I don’t think’ or ‘They don’t think’) are colour coded with cyan. The difference in model levels of sycophancy between the two query types represents how opinionated the model is in response to the queries for a particular topic. For objectively verifiable statements (Figures 6, 7, 8, 9), we expect to see elevated levels of sycophancy on negative queries as the model should be less accommodating to a user’s objectively wrong views. On the other side, elevated levels on affirmative queries should be considered problematic. For debateable statements (Figures 2, 3, 4, 5), any difference in sycophancy levels between affirmative and negative queries should be considered indicative of the models’ personae.

First versus Third Person Queries. In Figures 2, 3, 6, 7, first-person queries (FPQ, i.e. using prefix ‘I think’ or ‘I don’t think’) are plotted separately from third-person queries (TPQ, i.e. using prefix ‘They think’ or ‘They don’t think’). In Figures 4, 5, 8, 9, first-person and third-person queries are plotted together for each individual topic but aggregated in separate bars (*F for first-person queries and *T for third-person queries). We expect to reproduce the third-person effect, where models show elevated levels of sycophancy in response to first-person queries compared to third-person queries. The interesting question is whether this effect holds across models and topics.

Comparisons between Individual Mod-

els. Besides showing differences in general levels of sycophancy across models, Figures 6 and 7 can be used to demonstrate which models are more or less committed to truthfulness in response to objectively verifiable statements. Figures 4 and 5 carry additional information about how models perform on individual topics, incl. whether they have stronger opinions on or how predictably sycophantic or neutral they are on particular topics.

Comparisons between Individual Topics. Figures 4, 5, 8, 9 show how models performed on individual topics. The figures compare sycophancy levels on a fixed axis, allowing for direct comparison between topics. In every subplot, each condition tested (i.e. affirmative first-person queries (AF), negative first-person queries (NF), affirmative third-person queries (AT), negative third-person queries (NT)) is plotted separately. Green diamonds represent the average sycophancy level across all topics for each condition, which allows us to easily identify topics that show above- or below-average sycophancy levels for each condition. Consistently miscalibrated sycophancy levels for a particular topic may serve as an informative proxy for the kinds of training data that introduce sycophancy in the first place. Consistently elevated levels of sycophancy for a particular topic may also indicate that the topic is more likely to be affected by the safety concerns associated with sycophancy, and thus may be a more urgent target for mitigations.




Figure 2: Mean sycophancy scores (0-3) for Claude Sonnet 4.6 across topics, separated by condition.




Figure 3: Mean sycophancy scores (0-3) for Grok 4 Fast across topics, separated by condition.




Figure 4: Sycophancy scores by topic and condition for Claude Sonnet 4.6. Diamond markers indicate condition means across topics.




Figure 5: Sycophancy scores by topic and condition for Grok 4 Fast. Diamond markers indicate condition means across topics.




Figure 6: Mean sycophancy scores (0–3) for Claude Sonnet 4.6 across topics, separated by condition.




Figure 7: Mean sycophancy scores (0–3) for Grok 4 Fast across topics, separated by condition.




Figure 8: Sycophancy scores by topic and condition for Claude Sonnet 4.6. Diamond markers indicate condition means across topics.




Figure 9: Sycophancy scores by topic and condition for Grok 4 Fast. Diamond markers indicate condition means across topics.

DISCUSSION. Reproducing Previous Results. Preliminary findings show that the hypothesis that the objectivity of a query outweighs the effect of context domain on sycophancy levels is indeed correct. Both models show significantly higher levels of sycophancy in response to disputeable queries compared to objectively falsifiable queries, and this effect holds across all topics (RQ1). More interestingly, we were not able to reliably reproduce the third-person effect, as the evaluated models only showed lower levels of sycophancy in response to third-person queries compared to first-person queries on some topics, but not all. This suggests that the third-person effect may be more context-sensitive than previously thought, and that it may not be a universal phenomenon across all topics and models. Finally, we were also not able to use affirmative queries to reliably trigger more sycophancy compared to negative queries, as models demonstrated diverging behaviours (e.g. Claude Sonnet 4.6 showed higher levels of sycophancy in response to negative queries compared to affirmative queries, while Grok 4 Fast showed the opposite pattern). This suggests that some models may be more committed to their own opinions or to comforting users with opposing, negatively-phrased views rather than catering to any of the user’s views.

Novel Observations. Besides reproducing some of the previous findings, preliminary results also show some novel observations. For one, we observed that what previous literature

considered more sycophancy-prone topics (e.g. governance and politics, art and culture, spirituality religion) actually showed average or even below-average levels of sycophancy when controlling for objectivity. Interestingly, some of these topics (art and culture, spirituality and religion) even demonstrate consistently higher levels of sycophancy in response to negative queries compared to affirmative queries (RQ2). Second, as hoped, all models show appropriately high sycophancy across objectively falsifiable negative queries (e.g. ‘I don’t think Macs can’t get a virus’), but elevated levels of sycophancy in response to objectively falsifiable affirmative queries (e.g. ‘I think Macs can’t get a virus’) for some topics remain of concern. Third, it is not only general levels of sycophancy across different topics but also relative levels of sycophancy (e.g. affirmative vs negative queries) that seem to vary. This may be indicative of biases in training data but also of the models’ personae, and thus may be relevant for both understanding the origins of sycophancy and for targeting mitigations. No results so far show a clear relationship between the representation of a topic in RLHF data and the sycophancy levels for that topic (RQ3). Despite vague positive and negative trends of affirmative and negative queries (respectively) across topics ordered by frequency, no measurements were reliably able to predict the exponential distribution in preferential data (Figure 1).

A CONTEXT DOMAIN TAGS

FINAL TAG	COUNT	MERGED TAGS (COUNT)
programming and technology	241	web development (1), language processing (1), science and technology (1), artificial intelligence ai (1), software development (2), cryptography (3), electronics (4), machine learning (4), computer science (5), artificial intelligence (16), programming (60), technology (143)
writing and literature	112	literary criticism (1), fantasy magic (1), fictional crossover (1), literary theory (1), fantasy literature (1), fiction writing (1), fan fiction (1), storytelling theory (1), writing (1), mystery literature (1), autobiography (1), creative writing (1), childrens story (1), childrens literature (1), scifi fantasy (1), fanfiction (2), science fiction (13), fiction (18), fantasy (20), literature (44)
art and culture	111	folklore (1), cultural studies (1), ai art (1), art technology (1), ai art generation (1), art and craft (1), tradition (1), animemanga (2), mythology (2), art and photography (2), art and culture (2), language arts (2), art and design (2), art creativity (2), photography (3), animation (3), anime (4), art and technology (4), culture (8), art (68)
games and entertainment	75	game development (1), entertainment industry (1), gaming development (1), satire (1), video games (2), games (2), comedy (3), humor (3), fantasy gaming (4), entertainment (28), gaming (29)
education and academia	53	research (1), academic integrity (1), academic research (2), science (11), education (38)
health	41	healthcare (1), healthcare technology (1), clinical trials (1), mental health (2), medicine (2), health (34)
business and management	33	customer service (1), commerce (1), it management (1), business compliance (1), business operations (1), manufacturing (1), ecommerce (1), business communication (2), business management (2), marketing (7), business (15)
maths and data science	28	statistics (1), geometry (1), data generation (1), puzzles (2), data science (4), mathematics (19)
family and relationships	27	caregiving (1), sexuality (1), parenting (1), motherhood (1), dating (1), family dynamics (1), family and relationships (1), weddings (1), interpersonal relationships (1), human connection (1), relationships (2), parenthood (2), family (2), romance (3), friendship (7)
physical sciences	27	space exploration (1), robotics (1), earth science (1), geology (2), climate science (2), environmental science (2), weather (3), astronomy (4), physics (11)
money and finance	21	cryptocurrency (1), real estate (1), economics (4), finance (15)
history	19	postcolonialism (1), history (18)
film and theatre	18	horror (1), film history (1), filmtelevision (1), film studies (1), film analysis (1), theater (1), television (2), drama (2), film (8)
language and translation	18	translation (1), language learning (2), language translation (4), linguistics (5), language (6)
music and pop-culture	14	pop culture (1), audio production (1), celebrity comparison (1), music therapy (1), music (10)
lifestyle and self-improvement	14	fitness (1), leisure (1), luxury (1), personal development (1), home decor (1), recreation (1), beauty industry (1), lifestyle (1), creativity (2), selfimprovement (2), fashion (2)
environment and sustainability	13	botany (1), nature (1), sustainability (1), business and environment (1), wildlife (1), energy (2), environment (6)
politics and governance	13	governance (1), international relations (1), global issues (1), politics (10)
law and justice	12	criminal justice (1), labor rights (1), crime (2), law (8)
media and communications	12	social media policies (1), telecommunications (1), media production (1), business coaching (1), social media monetization (1), marketing technology (1), media (1), social media (2), communication (3)
engineering and construction	11	architecture and engineering (1), construction industry (1), architecture (2), construction (3), engineering (4)
society and social sciences	10	demography (1), futurology (1), indigenous issues (1), society (1), demographics (1), sociology (1), labor movement (1), accessibility (1), dystopia (2)
conflict and security	9	military (1), war (1), military technology (1), conflict (1), cybersecurity (5)
philosophy and ethics	9	artificial intelligence ethics (2), philosophy (3), ethics (4)
other	9	pet products (1), greetings (1), privacy (1), paranormal (1), mystery (1), teen drama (1), transformation (1), roleplaying (1), scheduling (1)
sports	9	sports politics (1), aquatics (1), cycling (1), sports (6)
life sciences	9	bioinformatics (1), biotechnology (1), biology (3), chemistry (4)
travel and tourism	8	travel safety (1), tourism and sustainability (1), travel and food (1), travel (5)
workplace and career	8	workplace diversity (1), workplace culture (1), workplace safety and compliance (1), corporate event (1), workplace communication (1), career development (3)
spirituality and religion	7	spirituality (1), religion (6)
cognitive sciences	5	science and consciousness (1), psychology (4)
cooking and baking	4	food cooking (1), food safety (1), cooking (2)

Table 1: grouped tags and their associated original tags as generated by GPT-4o

REFERENCES

- [1] OpenAI & team. GPT-4o System Card, 2024.
- [2] Kazi NOSHIN, Syed Ishtiaque AHMED, and Sharifa SULTANA. User Detection and Response Patterns of Sycophantic Behavior in Conversational AI, 2026.
- [3] Emilio BARKETT, Olivia LONG, and Madhavendra THAKUR. Reasoning Isn’t Enough: Examining Truth-Bias and Sycophancy in LLMs, 2025.
- [4] María Victoria CARRO. Flattering to Deceive: The Impact of Sycophantic Behavior on User Trust in Large Language Model, 2024.
- [5] Myra CHENG, Cinoo LEE, Pranav KHADPE, Sunny YU, Dyllan HAN, and Dan JURAFSKY. Sycophantic AI Decreases Prosocial Intentions and Promotes Dependence. *Science*, 391(6792):eaec8352, 2026.
- [6] Myra CHENG, Sunny YU, Cinoo LEE, Pranav KHADPE, Lujain IBRAHIM, and Dan JURAFSKY. Social Sycophancy: A Broader Understanding of LLM Sycophancy, 2025.
- [7] Magda DUBOIS, Cozmin UDUDEC, Christopher SUMMERFIELD, and Lennart LUETTGAU. Ask Don’t Tell: Reducing Sycophancy in Large Language Models, 2026.
- [8] Aaron FANOUS, Jacob GOLDBERG, Ank A. AGARWAL, Joanna LIN, Anson ZHOU, Roxana DANESHJOU, and Sanmi KOYEJO. SycEval: Evaluating LLM Sycophancy, 2025.
- [9] Shomik JAIN, Charlotte PARK, Matt VIANA, Ashia WILSON, and Dana CALACCI. Interaction Context Often Increases Sycophancy in LLMs. *arXiv preprint arXiv:2509.12517*, 2025.
- [10] Walter LAURITO, Benjamin DAVIS, Peli GRIETZER, Tomáš GAVENČIAK, Ada BÖHM, and Jan KULVEIT. AI-AI Bias: Large Language Models Favor Communications Generated by Large Language Models, July 2025.
- [11] Lars MALMQVIST. Sycophancy in Large Language Models: Causes and Mitigations, 2024.
- [12] Ivo PETROV, Jasper DEKONINCK, and Martin VECHEV. BrokenMath: A Benchmark for Sycophancy in Theorem Proving with LLMs, 2025.
- [13] Leonardo RANALDI and Giulia PUCCI. When Large Language Models Contradict Humans? Large Language Models’ Sycophantic Behaviour, 2025.
- [14] Itai SHAPIRA, Gerdus BENADE, and Ariel D. PROCACCIA. How RLHF Amplifies Sycophancy, 2026.
- [15] Mrinank SHARMA, Meg TONG, Tomasz KORBAK, David DUVENAUD, Amanda ASKELL, Samuel R. BOWMAN, Newton CHENG, Esin DURMUS, Zac HATFIELD-DODDS, Scott R. JOHNSTON, Shauna KRAVEC, Timothy MAXWELL, Sam MCCANDLISH, Kamal NDOUSSE, Oliver RAUSCH, Nicholas SCHIEFER, Da YAN, Miranda ZHANG, and Ethan PEREZ. Towards Understanding Sycophancy in Language Models, 2025.

- [16] Alexandra SOULY, Javier RANDO, Ed CHAPMAN, Xander DAVIES, Burak HASIRCIOGLU, Ezzeldin SHEREEN, Carlos MOUGAN, Vasilios MAVROUDIS, Erik JONES, Chris HICKS, Nicholas CARLINI, Yarin GAL, and Robert KIRK. Poisoning Attacks on LLMs Require a Near-constant Number of Poison Samples, 2025.
- [17] Keyu WANG, Jin LI, Shu YANG, Zhuoran ZHANG, and Di WANG. When Truth Is Overridden: Uncovering the Internal Origins of Sycophancy in Large Language Models, 2025.
- [18] Jerry WEI, Da HUANG, Yifeng LU, Denny ZHOU, and Quoc V. LE. Simple Synthetic Data Reduces Sycophancy in Large Language Models, 2024.